

IMPROVING MULTI-LABEL DEEP LEARNING MODELS FOR CLASSES WITH LONG-TAIL DISTRIBUTION

A PROJECT THESIS

Submitted in Partial fulfilment of the requirements for the award of the degree of

BACHELOR OF TECHNOLOGY IN ELECTRICAL AND ELECTRONICS ENGINEERING

by

PRASHANT MISHRA (Roll No. 521218)

RAKESH YAMJALA (Roll No. 521260)

K ABHISHEK RAJ (Roll No. 521143)

SANDEEP KUMAR AAZAD (Roll No. 521230)

Supervisor :

Dr. Sri Phani Krishna Karri



Department of Electrical Engineering

National Institute of Technology (NIT) Andhra Pradesh

April 2025

APPROVAL SHEET

This thesis entitled “Improving Multi-Label Deep Learning models for classes with Long-Tail Distribution” by “**Rakesh Yamjala, Prashant Mishra, K Abhishek Raj, Sandeep Kumar Aazad**” is approved for the degree Bachelor of Technology in Electrical and Electronics Engineering from National Institute of Technology Andhra Pradesh.

Supervisor

Dr. Sri Phani Krishna Karri
Assistant Professor
Department of Electrical Engineering
N.I.T Andhra Pradesh

Head of the Department

Dr. T. Ramesh
Assistant Professor
Department of Electrical Engineering
N.I.T. Andhra Pradesh

Date: _____

Place: _____

DECLARATION

We declare that this written submission represents our ideas in our own words and where others' ideas or words have been included, we have adequately cited and referenced the original sources. We also declare that we have adhered to all principles of academic honesty and integrity and have not misrepresented or fabricated or falsified any idea/data/fact/source in our submission. We understand that any violation of the above will be cause for disciplinary action by the Institute and can also evoke penal action from the sources which have thus not been properly cited or from whom proper permission has not been taken when needed.

(Signature)

Rakesh Yamjala

Roll No : 521260

Date : 28/04/2025

(Signature)

Prashant Mishra

Roll No: 521218

Date: 28/04/2025

(Signature)

Sandeep Kumar Aazad

Roll No : 521230

Date : 28/04/2025

(Signature)

K. Abhishek Raj

Roll No: 521143

Date: 28/04/2025

CERTIFICATE

It is certified that the work contained in the thesis titled “Improving Multi-Label Deep Learning models for classes with Long-Tail Distribution” by “**Rakesh Yamjala**, bearing Roll No: **521260**, **Prashant Mishra**, bearing Roll No: **521218**, **Sandeep Kumar Aazad**, bearing Roll No: **521230** and K. **Abhishek Raj**, bearing Roll No.: **521143**” has been carried out under my supervision and this work has not been submitted elsewhere for a degree.

Signature of the Supervisor
Dr. Sri Phani Krishna Karri
Assistant Professor
Dept. of Electrical Engineering
NIT Andhra Pradesh
April, 2025

Signature of the HOD
Dr. T. Ramesh
Assistant Professor
Dept. of Electrical Engineering
NIT Andhra Pradesh
April, 2025

ACKNOWLEDGEMENT

We extend our heartfelt gratitude to all those who have contributed to the successful completion of this project. Their unwavering support and guidance have been invaluable throughout this journey.

First and foremost, we express our sincere appreciation to our project guide, **Dr. Sri Phani Krishna Karri**, for his exceptional mentorship, valuable insights, and continuous support. His expertise has been instrumental in shaping the direction of this project.

We are also thankful to **Dr. Tejavathu Ramesh**, Head of the Department of Electrical and Electronics Engineering (EEE), for his encouragement and support throughout our academic endeavours.

We would like to extend our gratitude to Ph.D. scholar **R S Naga Kumar** for guidance and thesis writing, which have significantly contributed to the development of this project.

Finally, we pay our respects and love to our parents and friends for their encouragement and support.

Date: 28/04/2025

Place: National Institute of Technology Andhra Pradesh

Rakesh Yamjala (Roll No. 521260)

Prashant Mishra (Roll No. 521218)

K. Abhishek Raj (Roll No. 521143)

Sandeep Kumar Aazad (Roll No. 521230)

Department of Electrical Engineering
National Institute of Technology Andhra Pradesh,
Tadepalligudem – 534101

ABSTRACT

This project addresses the critical challenge of multi-label classification in medical imaging, focusing on the automated detection of co-occurring ocular diseases from retinal fundus photographs. Utilizing the ODIR-5K dataset—a benchmark repository of 5,000 paired left- and right-eye images annotated with up to eight labels (Normal, Diabetes, Glaucoma, Cataract, Age-related Macular Degeneration, Hypertension, Myopia, and Others) we systematically evaluate and enhance deep learning models under severe class imbalance. Four established convolutional neural network backbones (VGG16, InceptionV3, ResNet50, and ResNet101) are adapted to the multi-label setting by replacing the standard softmax layer with independent sigmoid outputs. To combat the “negative dominance” effect where the abundance of absent labels overwhelms positive signals, we implement a comprehensive suite of loss functions: unweighted and class-weighted Binary Cross-Entropy (BCE), Focal Loss, Hill Loss, Dice Loss, and hybrid formulations (BCE+Hill, Dice+Hill). A stratified five-fold cross-validation framework ensures that each disease’s long-tail distribution is faithfully represented during training and evaluation.

All models are trained with full fine-tuning of pretrained weights, followed by a two-stage protocol that freezes early convolutional blocks after five epochs and fine-tunes deeper layers with reduced dropout rates. The Adam optimizer (initial learning rate $1e-4$) promotes stable convergence. Among the architectures tested, ResNet101 achieves the highest performance 66 % test accuracy and a macro-averaged F1-score of 0.62 when optimized with the weighted Binary Cross-Entropy (BCE), demonstrating superior sensitivity to both common and rare labels. These results underscore the importance of backbone depth, tailored loss-function engineering, and robust training strategies for multi-label classification on long-tail medical datasets. Future work will explore few-shot and zero-shot vision language approaches to further enhance detection of underrepresented or novel conditions without extensive retraining.

Key Words: Multi-Label Classification, Long-Tail Distribution, Ocular Disease Diagnosis, ODIR-5K Dataset, Deep Convolutional Neural Networks, Loss Function Engineering, Few-Shot and Zero-Shot Learning

LIST OF FIGURES

Chapter	Figure No.	Title	Page No.
1	Fig 1.1	Sample fundus images from the ODIR-5K dataset showing variations across Normal, ARMD, Cataract, and Diabetic Retinopathy classes	2
2	Fig 2.1	Long-Tail Distribution of Diseases in the ODIR Dataset	9
3	Fig 3.1	Fundus images with single and multi-labels [Article by Lu Zhenzhen]	16
4	Fig 4.1	Architecture of ResNet50	23
4	Fig 4.2	Architecture of InceptionV3	24
4	Fig 4.3	Architecture of VGG16	25
4	Fig 4.4	Architecture of ResNet101	26
4	Fig 4.5	Comparison Graph of different Models	30
5	Fig 5.1	Performance comparison of different neural network architectures	36
6	Fig 6.1	Comparison of Architecture using Weighted BCE Loss	39

TABLE OF CONTENTS

Title	Page No.
APPROVAL SHEET	ii
DECLARATION	iii
CERTIFICATE	iv
ACKNOWLEDGEMENT	v
ABSTRACT	vi
LIST OF FIGURES	vii
LIST OF ABBREVIATIONS	ix
1.INTRODUCTION	1
1.1 The ODIR-5K Dataset	1
1.2 State of the Art	3
1.2.1 Convolutional Neural Network Backbones	3
1.2.2 Specialized Loss Functions for Long-Tailed Learning	3
1.2.3 Data Augmentation and Validation	4
1.3 Scope of the Work	4
1.4 Outline of the Chapters	5
2.REVIEW OF LITERATURE	6
2.1 General	6
2.2 Prior Research Works	6
2.2.1 Robust Asymmetric Loss for Multi-Label Long-Tailed Learning	7
2.2.2 Multi-Label Classification with Graph Convolution and Self-Supervised Learning	7
2.2.3 Ocular Disease Detection from Fundus Images	8
2.2.4 Utilizing ResNet-101 for Feature Extractions	9
2.2.5 Loss Function for Multi Label Classifications	11
2.2.6 Data Augmentation and Preprocessing	12
2.2.7 Handling class Imbalance	12
2.3 Conclusions	13
3.DATASET AND DATA PREPROCESSING	14
3.1 Introduction	14

3.2 Dataset Overview	15
3.3 Challenges	16
3.4 Data Augmentation	17
4.METHODOLOGY	19
4.1 General	19
4.2 Model Architectures Used	21
4.2.1 ResNet50	21
4.2.3 InceptionV3	23
4.2.5 VGG16	25
4.2.7 ResNet101	26
4.3 Loss Functions and Impact	27
4.4 Model Training Strategy	28
4.5 Evaluation Metrics	29
5.ADVANCED TECHNIQUES FOR ROBUSTNESS AND ACCURACY ENHANCEMENT	31
5.1 General	31
5.2 Literature Review	32
5.2.1 Binary Cross-Entropy (BCE) Loss	32
5.2.2 Weighted Binary Cross-Entropy (Weighted BCE) Loss	33
5.2.3 Hill Loss	33
5.2.4 Dice Loss	34
5.2.5 Hybrid BCE + Hill Loss	35
5.3 Conclusions	35
6.RESULTS AND CONCLUSION	38
6.1 Result	38
6.1.1 Overall Performance	38
6.1.2 Loss Function Comparison	39
6.1.3 Cross-Validation Stability	39
6.1.4 Summary of Key Results	39
6.2 Conclusion	40
6.2.1 Key Findings	40
6.2.2 Implications	40
6.2.3 Limitations	40
6.2.4 Future Work	40
REFERENCES	41
BIO-DATA	42

LIST OF ABBERRVIATIONS

Abbreviation	Full Form
CNN	Convolutional Neural Network
BCE	Binary Cross-Entropy
BCE + Hill	Hybrid Binary Cross-Entropy + Hill Loss
Dice Loss	Sørensen–Dice Coefficient Loss
ODIR-5K	Ocular Disease Intelligent Recognition 5K Dataset
AUC	Area Under Curve (ROC-AUC)
ROC	Receiver Operating Characteristic
GPU	Graphics Processing Unit
ReLU	Rectified Linear Unit
ROC-AUC	Receiver Operating Characteristic - Area Under Curve
SGD	Stochastic Gradient Descent
SimCLR	Simple Framework for Contrastive Learning of Visual Representations
VGG16	Visual Geometry Group 16-layer network
ResNet50	Residual Network with 50 Layers
ResNet101	Residual Network with 101 Layers
Hill Loss	Robust Margin-Based Confidence Loss
Adam	Adaptive Moment Estimation (Optimizer)
GPU	Graphics Processing Unit
ReLU	Rectified Linear Unit
ROC-AUC	Receiver Operating Characteristic - Area Under Curve
SGD	Stochastic Gradient Descent
Hill Loss	Robust Margin-Based Confidence Loss

CHAPTER 1

INTRODUCTION

1.1 General

In many real-world applications, an image or any data sample can belong to more than one category at the same time. This is known as multi-label classification, and it differs fundamentally from the more familiar single-label setting, where each input is assigned exactly one class. In the multi-label case, we must account for potentially complex relationships among labels: some diseases often co-occur, while others are mutually exclusive. We also face a very large and often sparse label space, meaning that at inference time we need to decide, for each label independently, whether it applies to a given image.

These requirements introduce several challenges. First, standard metrics like overall accuracy can be misleading; we need per-label precision, recall, and specialized measures (e.g., Hamming loss, Jaccard index) to fully capture model performance across both frequent and rare classes. Second, the design of loss functions becomes critical. If we simply use unweighted binary cross-entropy, the overwhelming number of negative labels in each sample will dominate training and cause the model to under-predict positive labels (*Distribution-Balanced Loss for Multi-Label Classification*, 2023). Addressing this “negative dominance” often requires weighted or hybrid loss formulations that amplify the signal from rare classes without destabilizing the gradient for head classes.

Moreover, multi-label models must produce well-calibrated probability estimates so that meaningful, per-class thresholds can be selected—an essential feature for clinical decision support where false positives and false negatives carry very different risks. Capturing inter-label dependencies (for example, modelling that diabetic retinopathy and hypertension often co-occur) can further improve performance, but adds complexity to both network design and training. Finally, practical deployment in a healthcare setting demands not only high predictive power but also interpretability, robustness to data shifts, and efficient inference, making multi-label classification in medical imaging a rich area of ongoing research.

1.1.1 The ODIR-5K Dataset

The Ocular Disease Intelligent Recognition (ODIR-5K) dataset comprises paired left and right eye fundus images from 5,000 patients, each annotated with up to eight possible labels: Normal, Diabetes, Glaucoma, Cataract, Age-related Macular Degeneration, Hypertension, Myopia, and Others (*ODIR-5K: Benchmarking and Challenges*, 2022). This dataset’s paired design enables exploration of bilateral disease patterns, although in this work each eye image is treated as an independent sample. A defining characteristic of ODIR-5K is its long-tail label distribution: while

common conditions such as Normal and Diabetes appear in over 60% of images, less prevalent diseases like Hypertension and Myopia may account for fewer than 5% of samples. This severe imbalance poses a significant obstacle to learning algorithms, which naturally gravitate toward head classes and risk failing to detect rare but clinically important diseases.




Fig 1.1 Sample fundus images from the ODIR-5K dataset showing variations across Normal, ARMD, Cataract, and Diabetic Retinopathy classes, highlighting the visual diversity crucial for model training.

1.1.2 Challenges in Long-Tail Multi-Label Learning

Two intertwined challenges define the long-tail, multi-label learning problem. First, class imbalance leads to disproportionate learning: models trained with unweighted loss favor majority labels, resulting in low recall for under-represented conditions. Second, negative dominance arises because, for any given image, absent labels vastly outnumber present ones. Unweighted binary cross-entropy loss thus becomes dominated by negative examples, further suppressing gradient

contributions from positive labels (*Advanced Focal and Contrastive Losses for Imbalanced Data, 2023*).

To build a clinically viable screening tool, it is essential to develop a training pipeline that simultaneously enhances sensitivity to tail-class instances, maintains overall performance on head classes, and provides reliable evaluation under imbalanced conditions. Moreover, the extreme sparsity of rare labels necessitates loss functions and sampling strategies that can amplify minority-class signals without destabilizing optimization for majority classes. Addressing these issues is critical to ensure that automated systems do not overlook less common but vision-threatening pathologies.

1.2 STATE OF THE ART

The current state of the art in long-tail multi-label classification integrates advances in CNN architectures, loss-function engineering, data augmentation, and efficient learning paradigms.

1.2.1 Convolutional Neural Network Backbones

Early deep architectures such as VGG16 characterized by a deep stack of 3×3 convolutional filters offered simplicity but incurred high parameter counts and limited representational capacity for subtle patterns (*Advances in VGG Variants for Medical Imaging, 2023*). InceptionV3 improved parameter efficiency by introducing parallel convolutional paths of varying kernel sizes, enabling the network to capture multi-scale features with fewer parameters (*Improved Inception Modules for Healthcare Applications, 2022*). ResNet50 and ResNet101 further advanced depth by employing residual connections, which mitigate vanishing gradients and facilitate the training of very deep networks (*Residual Network Optimizations for Imbalanced Classification, 2024*).

In multi-label settings, these backbones are adapted by replacing the final softmax layer with independent sigmoid activations, allowing each label to be predicted concurrently. Transfer learning from ImageNet-pretrained weights remains the standard approach in medical imaging due to the scarcity of labelled data, providing a strong initialization and speeding convergence.

1.2.2 Loss Function Innovations

To address class imbalance and negative dominance, numerous specialized loss functions have been proposed. **Weighted binary cross-entropy** scales each class's loss by an inverse frequency factor, emphasizing rare labels (Class-Balanced Loss for Sparse Multi-Label Data, 2022). **Focal loss** down-weights easy negatives and focuses training on hard positives, mitigating the overwhelming effect of abundant negative examples (Enhanced Focal Loss Variants for Multi-Label Learning, 2023). **Hill loss** shapes the loss landscape to privilege false negatives in a "hill" form, balancing head and tail classes (Robust Hill Loss for Long-Tail Learning, 2023). **Dice loss**, originally developed for segmentation, directly maximizes overlap metrics and alleviates class

imbalance (Review of Dice and Overlap Loss in Medical Imaging, 2024). **Distribution-balanced loss** integrates label co-occurrence statistics with negative-tolerant regularization to further improve long-tail performance.

Hybrid combinations such as binary cross-entropy combined with Hill loss or Dice loss have demonstrated the ability to leverage complementary strengths of different formulations, yielding more robust optimization in multi-label scenarios.

1.2.3 Data Augmentation and Validation

Basic augmentation techniques horizontal and vertical flips, small rotations, and zoom have been shown to increase the effective sample size of minority classes and improve generalization (Comprehensive Survey of Data Augmentation for Medical Multi-Label Imaging, 2023). More advanced strategies like Mix-up, Cut-Mix, and GAN-based synthetic generation remain underexplored in medical multi-label contexts but hold promise for further balancing class distributions. To ensure robust evaluation on imbalanced data, **stratified k-fold cross-validation** preserves the long-tail distribution across folds, providing reliable estimates of model generalization and reducing overfitting risk.

1.3 SCOPE OF THE WORK

This thesis undertakes a systematic exploration of multi-label classification techniques tailored for long-tailed medical imaging data. The specific objectives are:

1. **Multi-label model development**

Design and implement a deep neural network for multi-label fundus disease classification using four standard CNN backbones VGG16, InceptionV3, ResNet50, and ResNet101 with independent sigmoid outputs. Establish baseline performance using unweighted binary cross-entropy loss.

2. **Loss function evaluation**

Assess a suite of specialized loss functions weighted binary cross-entropy, focal loss, Hill loss, Dice loss, and hybrid combinations and quantify their impact on sensitivity to under-represented classes without sacrificing overall accuracy

3. **Efficient learning paradigms**

Incorporate data-efficient techniques such as few-shot learning (e.g., prototype-based frameworks) and zero-shot vision-language methods to mitigate limited labelled data scenarios, enabling detection of rare or novel conditions.

4. **Training strategy under stratified cross-validation**

Adopt a two-stage training regimen: first, train all network layers end-to-end; then, freeze early convolutional blocks and fine-tune deeper layers with reduced dropout rates to improve generalization.

5. Comprehensive evaluation

Apply rigorous evaluation metrics multi-label accuracy, per-class precision and recall, ROC AUC, and confusion analysis across stratified five-fold splits. Perform threshold tuning on validation sets to optimize detection performance for each label.

6. Future directions

Identify promising avenues for further improvement, including advanced augmentations, ensembling heterogeneous backbones, and exploration of transformer-based architectures for multi-label ocular diagnosis.

1.4 OUTLINE OF THE CHAPTERS**Chapter I: Introduction**

Defines multi-label long-tail challenges; presents motivating clinical application; outlines objectives and methodological overview.

Chapter II: Literature Review

Surveys multi-label classification fundamentals, imbalance-handling losses, data augmentation strategies, and efficient learning paradigms in medical imaging.

Chapter III: Materials and Methods

Describes ODIR-5K dataset specifics, preprocessing workflows, CNN architectures, loss functions, training protocols, and evaluation metrics.

Chapter IV: Experimental Results and Discussion

Reports quantitative and qualitative outcomes for baseline and enhanced models; analyzes impact on common and rare labels; presents confusion analyses.

Chapter V: Advanced Techniques and Ablation Studies

Examines few-shot and zero-shot learning integration, hybrid loss ablations, and fine-tuning procedures to isolate contributions of each component.

Chapter VI: Conclusions and Future Work

Summarizes findings, discusses limitations, and recommends directions for further research such as advanced augmentation, ensemble, and transformer models.

CHAPTER – 2

REVIEW OF LITERATURE

2.1 GENERAL

In the previous chapter, the key concepts related to multi-label classification models, long-tailed data distributions, and the specific challenges faced in medical image analysis were discussed. In this chapter, an extensive literature review is presented to build a comprehensive understanding of existing research efforts in this domain.

This chapter presents a comprehensive review of the existing literature related to multi-label classification models, emphasizing techniques for addressing long-tailed distributions in medical imaging datasets. The objective is to establish a solid foundation for the proposed improvements undertaken in this project.

Tackling the long-tailed distribution of data was identified as one of the central challenges during the project. Important research papers, such as "*Robust Asymmetric Loss for Multi-Label Long-Tailed Learning*" (2023) [1] and "*Multi-Label Classification of Fundus Images With Graph Convolutional Network and Self-Supervised Learning*" (2021) [2], highlight the severity of class imbalance issues and laid the groundwork for our approach. These studies not only revealed the limitations of conventional methods when dealing with imbalanced medical datasets but also motivated the adoption of novel loss functions and self-supervised learning techniques in our project to mitigate these challenges.

This literature review is primarily focused on a steady-state analysis of existing models in the context of multi-label ocular disease classification. Particular attention has been given to understanding the types of datasets employed, the architectural choices made, and the performance metrics achieved in previous works. By critically analyzing these aspects, we have sought to identify gaps and opportunities for further improvement. Building upon these insights, our project aims to enhance model accuracy and precision, particularly in recognizing rare disease classes, thereby contributing to more reliable and generalizable medical diagnostic systems.

2.2 LITERATURE REVIEW

Extensive literature survey has been conducted for this project work and is divided into eight parts subject-wise. These are:

2.2.1 Multi-Label Classification Techniques

The phenomenon of multi-label classification in medical imaging was first comprehensively addressed in the study "*Multi-Label Classification of Fundus Images with Graph Convolutional Network and Self-Supervised Learning*" (2021) [2]. In this work, the authors emphasized that the class distribution of ocular disease datasets exhibits a pronounced long-tailed shape, where the prevalence of different diseases varies significantly. Moreover, it is common for fundus images from symptomatic patients to exhibit multiple concurrent pathological labels, complicating the classification task.

The study proposed the use of Graph Convolutional Networks (GCNs) to model label correlations explicitly, recognizing that traditional methods often fail to capture interdependencies among disease classes. Additionally, self-supervised learning techniques were incorporated to pretrain feature extractors, leveraging the vast amount of unlabeled data available to improve feature robustness. Although the model achieved notable improvements in handling multi-label classification under class-imbalanced conditions, it was observed that further optimization was necessary to fully capture rare disease instances and minimize performance degradation on minority classes.

The conclusions drawn from this work highlighted the importance of modeling label dependencies and utilizing unlabeled data through self-supervised learning to enhance classification performance, especially in long-tailed, imbalanced medical datasets. These insights provided a crucial starting point for the present project, influencing the decision to adopt hybrid loss functions and further refine self-supervised feature extraction methods for improved accuracy and generalization.

2.2.2 Challenges of Long-Tail Distribution in Deep Learning

The challenges of long-tail distribution in deep learning for multi-label classification are notably addressed in the paper *Robust Asymmetric Loss for Multi-Label Long-Tailed Learning* (2023) [1]. In multi-label classification, a model predicts multiple labels for a single input, which is especially prevalent in fields like medical image recognition. In such domains, numerous studies have explored multi-label classification techniques, but they often fail to adequately address the issue of long-tailed distributions in the dataset.

A long-tailed distribution occurs when a small number of classes (head labels) are well-represented with abundant data, while a large number of other classes (tail labels) have much fewer data samples. This imbalance causes significant challenges in model performance. For tail labels, which

are underrepresented, the model tends to perform poorly due to the lack of sufficient training examples. On the other hand, the model becomes biased towards head labels, as these have a higher number of samples, leading to overfitting and poor generalization to tail classes.

This issue becomes particularly important in real-world applications, such as medical image classification, where the goal is not only to recognize common conditions (head labels) but also to identify rare diseases (tail labels). Addressing this imbalance is crucial to ensure that the model can generalize well to all classes, not just the dominant ones.

Several studies have attempted to solve this problem by re-sampling or re-weighting the training data. Re-sampling techniques either oversample the minority (tail) classes or under sample the majority (head) classes, while re-weighting assigns higher importance to the tail labels during training. These methods help mitigate the bias towards head labels, allowing the model to focus more on learning from tail labels. However, these approaches have primarily been explored in single-label classification settings, and their application in multi-label classification environments remains under-explored.

In traditional multi-label classification, Binary Cross-Entropy (BCE) Loss is commonly used to measure the difference between predicted probabilities and true labels. The BCE Loss for a single label can be formulated as:

$$\text{BCE}_{\text{loss}} = \frac{-1}{N} \sum_{i=1}^N N(y_i \log(p_i) + (1 - y_i) \log(1 - p_i))$$

where N is the number of classes, y_i is the true binary label (0 or 1) for the i -th class, and p_i is the predicted probability for the i -th class. The loss function penalizes large deviations between the predicted probabilities and the true labels, encouraging the model to improve its predictions across all labels. However, while BCE Loss works for balanced datasets, it does not inherently address class imbalance, especially in long-tailed distributions.

To solve the problem of long-tailed data in multi-label classification, more advanced loss functions, such as the Asymmetric Loss introduced in the paper, are being developed. These methods specifically account for the imbalance in label distributions and seek to give more weight to tail labels, improving model performance in such settings.

2.2.3 Ocular Disease Detection from Fundus Images

The Ocular Disease Intelligent Recognition (ODIR) dataset is used as the primary dataset in this study. It contains ophthalmic records of 5,000 patients, including age, bilateral fundus images, and diagnostic labels. The data were collected from multiple hospitals across China using various imaging devices, leading to diverse image qualities. Each patient is categorized into one or more

of eight disease classes, such as Diabetes, Glaucoma, and Cataract. The dataset captures a real-world clinical environment, reflecting practical diagnostic challenges. A major characteristic of ODIR is its **long-tailed distribution**, where common diseases dominate and rare diseases are underrepresented. This imbalance makes it difficult for models to generalize well across all disease classes.

Without proper handling, models tend to favor frequent classes while misclassifying rare ones. Thus, the ODIR dataset provides a valuable platform to explore techniques for managing imbalanced, multi-label medical data. It emphasizes the need for robust learning strategies to improve performance on underrepresented diseases. A more detailed discussion of the ODIR dataset and its properties is provided in the following chapter.




Fig. 2.1 Long-Tail Distribution of Diseases in the ODIR Dataset.

The figure illustrates the highly imbalanced distribution of disease categories within the ODIR dataset. Common conditions such as Normal, Diabetes, and Cataract are represented with a disproportionately higher number of instances compared to rarer conditions like Pathological Myopia and Other Abnormalities. This imbalance exemplifies the long-tailed distribution problem in multi-label medical image classification tasks, where models are at risk of being biased towards frequently occurring classes while underperforming on underrepresented, yet clinically significant, conditions.

2.2.4 Utilizing ResNet-101 for Feature Extraction

In the initial phase of this research, an extensive review of existing literature was undertaken to guide the selection of suitable deep learning architectures for multi-label classification on long-tailed datasets. Our first experimental model was **Inception V3**, implemented based on the paper "*Rethinking the Inception Architecture for Computer Vision*" (2015) [4]. Despite its sophisticated multi-scale processing capabilities, the model achieved a modest **Binary Cross-Entropy (BCE) accuracy of 54%**. The relatively lower performance was attributed to its difficulty in handling class imbalance and the complexity of fundus images.

Seeking to improve upon these results, we transitioned to **VGG16**, referring to "*Very Deep Convolutional Networks for Large-Scale Image Recognition*" (2014) [5]. Known for its simplicity and uniform architecture with small convolutional filters, VGG16 produced a **BCE accuracy of 55%**, showing a slight improvement but still insufficient for our clinical-grade application requirements.

Recognizing the need for a more robust architecture capable of addressing the vanishing gradient problem and enabling deeper network training, we adopted **ResNet-50**, inspired by "*Deep Residual Learning for Image Recognition*" (2015) [3]. Through the use of residual connections, **ResNet-50** facilitated better gradient flow and achieved a **BCE accuracy of 59%**, representing a significant leap in performance compared to earlier models.

In an effort to further optimize results, we finally implemented **ResNet-101**, a deeper variant of ResNet architectures, guided by the work "*A Novel ResNet101 Model Based on Dense Dilated Convolution for Image Classification*" (2022) [6]. By incorporating a greater depth and improved feature extraction through dilated convolutions, **ResNet-101** achieved the highest performance among all tested models, reaching a **BCE accuracy of 66%**. This significant improvement validated the importance of deeper networks for complex multi-label medical image classification tasks.

The superior performance of **ResNet-101** can be attributed to several key factors. Its deeper **residual architecture**, with 101 layers and identity (skip) connections, allows for efficient gradient flow, which is essential for learning fine-grained features needed to identify subtle retinal pathologies. Additionally, the **larger effective receptive field** afforded by the deeper stacking of layers enables the model to better distinguish small lesions from varied backgrounds, crucial in medical image classification. Finally, **ResNet-101** demonstrates superior generalization on long-tail classes. The increased capacity of the model, combined with the skip-connections, helps it effectively learn both common and rare classes without overfitting, thus addressing the imbalances present in the dataset.

Following the model experimentation, attention was directed toward addressing the persistent issue of class imbalance, a characteristic feature of the ODIR dataset. To further enhance the model's ability to generalize across rare and common disease classes, different loss functions were explored and optimized in subsequent stages of the study

2.2.5 Loss Functions for Multi-Label Classification

In the pursuit of optimizing multi-label classification on the ODIR dataset, a comprehensive evaluation of different loss functions was conducted, each offering distinct advantages and limitations. Initially, Binary Cross-Entropy (BCE) loss was employed as a baseline due to its simplicity and widespread use. While it provided a reasonable starting point, its uniform treatment of all classes proved inadequate in the face of severe class imbalance, often leading the model to prioritize majority classes and underperform on rarer labels.

To mitigate this imbalance, Weighted BCE was introduced, assigning higher importance to underrepresented classes. This adjustment resulted in a notable improvement in classification accuracy for minority labels, particularly enhancing sensitivity without significantly sacrificing overall performance. However, it still lacked the ability to directly optimize for multi-label set-level similarity.

To address this, Dice Loss was incorporated, which measures the overlap between predicted and true labels. Its design inherently favors rare class detection and proved effective in improving recall rates across minority classes. Nevertheless, models trained solely with Dice Loss exhibited unstable convergence patterns and required careful learning rate management.

Hill Loss was also explored, offering an alternative perspective by adjusting the margin constraints based on label frequency. It demonstrated a stronger balance between precision and recall, particularly benefiting classes with few examples. However, Hill Loss was found to be sensitive to hyperparameter settings, making its deployment more complex.

Recognizing the individual strengths of these loss functions, hybrid approaches were then evaluated. The Hybrid BCE+Hill Loss and Hybrid Dice+Hill Loss formulations combined stability with improved minority class sensitivity. These hybrid losses consistently outperformed single-loss functions, achieving the highest accuracies — particularly the Hybrid Dice+Hill Loss, which offered a more nuanced optimization of both head and tail classes. By blending the local pixel-wise accuracy of BCE and Dice with the distributional fairness promoted by Hill Loss, these hybrid strategies significantly enhanced the model's generalization ability across imbalanced multi-label categories.

Overall, the experimental findings underscore that no single loss function is universally optimal; instead, thoughtful combinations tailored to the data distribution yield the best performance in complex, long-tailed multi-label classification problems.

2.2.6 Data Augmentation and Preprocessing in Medical Image Analysis

Effective preprocessing and stratified data augmentation were integral to enhancing model performance for multi-label classification of medical fundus images. The preprocessing pipeline began with image resizing to a fixed resolution to ensure compatibility across different architectures, followed by pixel intensity normalization to standardize input distributions. Histogram equalization and adaptive contrast stretching techniques were selectively applied to improve the visibility of fine-grained pathological features, particularly beneficial for detecting early-stage anomalies.

To address the inherent class imbalance and small dataset size, stratified data augmentation strategies were implemented during training. These augmentations ensured that each class was well-represented during the augmentation process, maintaining the balance of the dataset. Conventional augmentations included random rotations (up to ± 20 degrees), horizontal and vertical flipping, slight zooming, and random shifts, all designed to simulate variability in real-world imaging scenarios while preserving anatomical realism. Additionally, photometric transformations such as random brightness, contrast, and gamma corrections were applied to model variations in imaging equipment and acquisition conditions. Careful constraints were placed on augmentation parameters to avoid introducing artifacts that could mislead the model. Together, the preprocessing and stratified augmentation pipeline significantly improved model robustness, reduced overfitting, and contributed to higher predictive performance, particularly in rare and complex diagnostic categories.

2.2.7 Handling Class Imbalance in Medical Datasets

Class imbalance presents a significant challenge in medical image analysis, particularly in multi-label classification tasks where certain pathological conditions are rare. In the ODIR dataset, common diseases such as diabetic retinopathy and glaucoma dominated the label distribution, while rare conditions like hypertension retinopathy and pathological myopia were severely underrepresented. To address this, targeted data-level strategies were employed. Augmentation techniques were selectively intensified for minority classes, including random rotations, shifts, and brightness variations, thereby synthetically increasing the diversity and quantity of rare samples without altering their clinical validity. This oversampling approach ensured that the model was exposed to a more balanced representation of disease categories during training, reducing the bias toward majority classes.

On the algorithmic side, specialized loss functions were leveraged to further counteract imbalance effects. Class-weighted Binary Cross-Entropy was initially implemented, assigning higher loss weights to minority class samples to make misclassifications more consequential during optimization. Additionally, hybrid loss formulations, combining the strengths of Weighted BCE and Dice Loss, were explored to jointly improve training stability and rare class sensitivity. These hybrid strategies allowed the model to maintain high precision on common labels while substantially boosting recall on underrepresented categories. Through the integration of data-level augmentation and loss-level rebalancing, the model achieved a more equitable performance distribution, which is essential for reliable and clinically meaningful diagnostic outputs.

2.3 CONCLUSIONS

In the preceding chapter, critical concepts surrounding multi-label classification models, long-tailed data distributions, and the inherent challenges within medical image analysis were discussed. Building on this foundation, the current chapter presented a comprehensive literature review aimed at deepening the understanding of existing research efforts related to these topics. Particular emphasis was placed on exploring techniques to address the issue of long-tailed distributions in medical imaging datasets, thereby establishing a strong groundwork for the methodological improvements proposed in this project.

Through an extensive survey of the literature, it became evident that tackling data imbalance is a pivotal challenge in multi-label medical image classification. Key studies, such as "Robust Asymmetric Loss for Multi-Label Long-Tailed Learning" and "Multi-Label Classification of Fundus Images with Graph Convolutional Network and Self-Supervised Learning," provided valuable insights into the limitations of traditional approaches and highlighted innovative solutions such as asymmetric loss functions and the use of self-supervised learning. In addition, a critical analysis of architectural choices revealed that deeper networks, particularly ResNet-101, offer substantial advantages in feature extraction and generalization, especially when dealing with rare disease classes in long-tailed datasets like ODIR.

Drawing upon the insights from these studies, this project strategically adopted advanced loss functions and deeper convolutional architectures to enhance model performance across both common and rare classes. By carefully addressing class imbalance and emphasizing robust feature extraction, the project aims to develop a more reliable and generalizable diagnostic system for ocular disease detection. This targeted approach ensures that the resulting model not only achieves high accuracy but also maintains sensitivity across underrepresented disease categories, contributing meaningfully to the broader field of automated medical diagnosis.

CHAPTER 3

DATASET AND DATA PREPROCESSING

3.1 INTRODUCTION

In the past decade, artificial intelligence (AI) has shown immense potential in transforming the healthcare industry, particularly in the field of medical imaging. One critical domain within this field is ophthalmology, where timely and accurate diagnosis of ocular diseases can prevent vision loss and significantly improve patients' quality of life. The application of deep learning, a subset of AI, to ophthalmic imaging has attracted considerable research interest. However, the success of deep learning models is heavily dependent on the availability of high-quality, annotated datasets that can effectively represent real-world challenges. The Ocular Disease Intelligent Recognition (ODIR-5K) dataset is one such significant contribution to the ophthalmology and AI communities. Released as part of the Grand Challenge associated with the 2019 International Symposium on Biomedical Imaging (ISBI), ODIR-5K provides a comprehensive collection of fundus images labeled for multiple ocular diseases. Its objective is to facilitate the development of machine learning algorithms capable of performing multi-label disease classification from retinal images, a task that closely mirrors actual clinical scenarios where a patient may suffer from more than one eye condition simultaneously.

Fundus photography captures the back of the eye, including the retina, optic disc, macula, and posterior pole. Many systemic and ocular diseases manifest detectable signs in these regions, making fundus imaging an essential diagnostic tool. By automating the analysis of these images, deep learning models could assist ophthalmologists in clinical decision-making, improving screening efficiency and reducing diagnostic errors. Unlike many existing datasets that focus on a single disease, ODIR-5K emphasizes multi-disease recognition, which reflects the real-world complexity of ophthalmic diagnostics. Each patient in the dataset is associated with a pair of images — one for each eye — and a set of diagnostic labels. These labels cover a broad spectrum of conditions, including diabetic retinopathy, glaucoma, cataracts, and age-related macular degeneration, among others. The dataset also includes demographic information such as age and sex, adding another layer of clinical relevance.

The ODIR-5K dataset stands out not only for its diversity but also for its challenges. Variations in image quality, resolution, and field of view, alongside class imbalance and the presence of noisy labels, provide a realistic environment for training and testing AI models. These challenges necessitate the use of sophisticated preprocessing techniques, robust model architectures, and careful evaluation metrics.

3.2 DATASET OVERVIEW

The ODIR-5K dataset comprises medical-grade fundus images from 5,000 patients, collected across various hospitals and eye clinics in China. These images are stored in JPEG format, organized along with a structured CSV file containing metadata and disease labels. Each record corresponds to a single patient and includes two images — one for the left eye and one for the right eye.

Each patient record contains the following fields:

- **ID:** A unique identifier.
- **Left-Fundus Image Path:** The file path to the left eye's fundus image.
- **Right-Fundus Image Path:** The file path to the right eye's fundus image.
- **Age:** The patient's age at the time of imaging.
- **Sex:** The patient's gender, recorded as male, female, or unspecified.
- **Left-Diagnostic Keywords:** Diseases or abnormalities diagnosed in the left eye.
- **Right-Diagnostic Keywords:** Diseases or abnormalities diagnosed in the right eye.
- **Labels:** Binary indicators for nine categories (Normal, Diabetes, Glaucoma, Cataract, Age-related Macular Degeneration, Hypertension, Myopia, Other Diseases, and Unspecified).

The disease labels are multi-hot encoded, meaning each patient may be labelled with multiple diseases simultaneously. This reflects the reality that patients often present with comorbid conditions rather than isolated pathologies.

The images vary in size, color balance, illumination, and field of view, as they were captured using different types of fundus cameras in varying clinical environments. Some images are high resolution and well-centered on the optic disc and macula, while others may suffer from blurring, poor focus, underexposure, or overexposure.




Fig.3.1 Fundus images with single and multi-labels. [Article by Lu Zhenzhen]

3.3 CHALLENGES

Working with the ODIR-5K dataset presents several challenges that closely mimic those faced in real-world clinical practice. These challenges must be addressed carefully to ensure robust model performance.

- **Multi-Label Classification:** One of the primary challenges is the dataset's inherently multi-label nature. Unlike traditional classification tasks where each sample belongs to only one class, each patient in ODIR-5K may exhibit signs of multiple ocular diseases. This complexity requires models to output a vector of probabilities rather than a single prediction, necessitating changes in loss functions (e.g., binary cross-entropy instead of categorical cross-entropy) and evaluation metrics (e.g., AUC, precision-recall, F1 score).
- **Class Imbalance:** Another significant challenge is class imbalance. Diseases like Normal and Diabetes have a large number of samples, whereas conditions like Myopia and Hypertension are underrepresented. Without proper handling, models tend to be biased towards predicting the majority classes while failing to identify rarer diseases. Techniques such as class weighting, data augmentation, SMOTE (Synthetic Minority Over-sampling Technique), and hybrid loss functions are often employed to address this issue.
- **Image Quality Variations:** The images were captured using different devices under various conditions, resulting in considerable heterogeneity in image quality. Issues such as Blurred images, Overexposed or underexposed images, Non-centered optic discs, Partial

visibility of the retina, complicate the feature extraction process. Robust preprocessing pipelines, including image enhancement, normalization, and data augmentation, are critical to mitigate these issues.

- **Metadata Noise:** While metadata such as age and sex could provide valuable auxiliary information for diagnosis, inconsistencies and missing entries limit their usefulness. Furthermore, textual diagnostic keywords can sometimes be ambiguous or incomplete, requiring careful parsing and cleaning before use.
- **Label Noise:** Some labels may be incorrect or incomplete, as they were assigned based on clinical reports, which may themselves contain errors or omissions. Label noise can misguide the training process, leading to degraded model performance. Techniques like label smoothing, semi-supervised learning, and consistency regularization are often used to counteract label noise.
- **Generalization Challenges:** Given that the data originates from a single geographic region, models trained solely on ODIR-5K may not generalize well to diverse global populations without further fine-tuning or external validation datasets. Bias in demographics, disease presentation, and imaging modalities must be considered when deploying such models in broader settings.

3.4 DATA AUGMENTATION

One of the most critical challenges in ODIR-5K is the severe imbalance between common and rare disease classes. Diseases such as diabetic retinopathy and normal cases dominate the dataset, whereas conditions like hypertension and myopia are comparatively rare. This imbalance causes models to become biased toward predicting the majority classes, resulting in poor sensitivity for minority classes.

Data augmentation can directly tackle this issue by **oversampling** the minority classes through extensive augmentation. Specifically:

- **Targeted Augmentation:** Applying heavier augmentations on underrepresented classes can synthetically increase their presence in the dataset without introducing label noise.
- **Augmentation Multiplication:** Minority class images can be augmented multiple times with different transformations to match the frequency of majority class images, thereby creating a more balanced training set.

Ocular diseases can manifest anywhere within the retinal image. For instance, diabetic retinopathy may cause lesions in various parts of the retina. Therefore, it is essential for the model to be spatially invariant — recognizing diseases regardless of their exact location within the image.

To enhance spatial invariance, we apply:

- **Random Rotations:** Rotating images by small random angles to prevent the model from learning fixed positional biases.
- **Horizontal and Vertical Flipping:** Reflecting images to teach the model that disease features can appear symmetrically in either eye.
- **Random Cropping and Zooming:** Encouraging the model to focus on partial regions, thus improving its ability to detect localized pathological signs.

Through these transformations, the model learns to recognize diseases based on structural features rather than their fixed spatial locations.

CHAPTER 4

MODEL DEVELOPMENT AND PERFORMANCE ANALYSIS

This chapter details the design, training, and evaluation of various deep learning models for multi-label classification using the ODIR-5K dataset. It includes the model architectures selected, loss functions experimented with, training methodology, and performance outcomes. Each aspect of the development pipeline is analysed in depth, supported by metrics, visualizations, and comparative results. The goal was to establish a reliable and scalable pipeline for ocular disease recognition, and this chapter outlines all the technical steps, insights, and iterations undertaken in the process.

4.1 General

The primary objective of this research was to develop a high-performing and robust multi-label image classification model aimed at the recognition of ocular diseases using retinal fundus images. The task diverges significantly from standard single-label classification problems, primarily due to the inherent complexities of multi-label annotation, where multiple disease conditions may coexist within a single retinal scan. These complexities introduce a series of challenges that impact both model design and evaluation strategies.

This study made use of the ODIR-5K (Ocular Disease Intelligent Recognition) dataset, a benchmark dataset designed specifically for ocular disease detection. The dataset includes 5,000 patient cases, each accompanied by retinal fundus images from both eyes, demographic information (age, sex), and expert-verified diagnostic labels that cover eight distinct disease categories, such as Diabetic Retinopathy, Cataract, Glaucoma, and Hypertension. Importantly, each image in the dataset may be labelled with more than one disease, making it a multi-label classification problem rather than a mutually exclusive single-label one.

In traditional single-label classification, models are trained to predict the most likely class from a set of disjoint classes. In contrast, multi-label classification requires the model to independently predict the presence or absence of each class, treating each disease label as a binary prediction task. This significantly increases the complexity of model training, especially in the context of medical imaging, where subtle pathological differences often exist between classes and diseases frequently co-occur.

Furthermore, the problem is exacerbated by the presence of class imbalance, where certain disease categories are overrepresented while others are rare. This imbalance leads to biased learning,

where the model tends to favor the majority classes, often at the cost of reduced recall for the minority ones. To counter this issue, a range of weighted and hybrid loss functions were explored during the study (discussed in later sections), each specifically designed to improve the model's sensitivity to underrepresented diseases.

Another technical barrier in ocular disease classification is the visual similarity between different disease types. Many diseases manifest with overlapping features such as similar patterns in the optic disc, vessel abnormalities, or lesions. Differentiating these nuanced visual cues requires models with strong discriminative capabilities and the ability to focus on fine-grained details—this influenced the selection of deeper architectures like ResNet101, which was eventually found to be the best performing model in this study.

To address these challenges, the research work involved building an end-to-end deep learning pipeline encompassing data preprocessing, architecture selection, loss experimentation, training strategies, and performance evaluation. Each stage of the pipeline was methodically optimized to ensure that the final model achieved maximum classification accuracy while remaining robust and generalizable across different disease categories.

Key innovations in the pipeline included:

- The use of stratified K-fold cross-validation to ensure robust model evaluation across the highly imbalanced dataset.
- Incorporation of basic data augmentation techniques to prevent overfitting and improve generalization.
- Exploration of hybrid loss functions to tackle multi-label class imbalance.
- Training deep models from scratch and using progressive layer unfreezing for effective feature learning.
- Careful tuning of dropout rates and optimizer hyperparameters.

The eventual aim was not only to maximize predictive performance, but also to ensure that the model could be trusted in a real-world clinical setting, where false negatives can have serious consequences. The outputs were designed to reflect class-wise probabilities rather than hard assignments, using sigmoid activation functions to support the multi-label nature of the task. A thorough performance evaluation using metrics tailored to multi-label classification, such as thresholded accuracy and AUC, was also conducted.

In summary, this section lays the foundation for the methodology discussed in detail in the subsequent chapters. The model workflow was designed to be modular, interpretable, and reproducible, allowing for future integration with more advanced techniques like transformer-based encoders, attention mechanisms, or threshold optimization algorithms that could further enhance clinical performance.

4.2 Model Architectures Used

A comprehensive review of existing literature was conducted to understand the strengths and limitations of various deep learning architectures, particularly in the context of long-tailed multi-label classification. This thorough examination allowed us to make informed decisions regarding model selection, ensuring that the chosen architectures align with the specific demands of our dataset and research objectives. After evaluating several studies and empirical results, we identified a set of promising models for experimentation, namely ResNet50, InceptionV3, VGG16, and ResNet101.

Each of these models has been extensively studied and applied in various domains, demonstrating their capability to extract rich and discriminative feature representations. To better appreciate their potential in addressing long-tailed distributions, we provide a detailed overview of the architectural design and operational principles of each model below. Understanding these foundational characteristics is critical to contextualizing their performance in the experiments that follow.

4.2.1 ResNet50

ResNet50 has emerged as one of the most effective architectures for addressing the challenges posed by long-tailed distributions in multi-label classification tasks. Its core innovation the introduction of residual learning directly contributes to its robustness in imbalanced data scenarios. In traditional deep networks, as the number of layers increases, models often encounter the problem of vanishing gradients, leading to degradation in performance. This issue is particularly harmful when learning from minority classes, which are already underrepresented and harder to classify.

Residual connections in ResNet50 mitigate this challenge by allowing the network to learn identity mappings more easily, ensuring that useful features from earlier layers are preserved throughout the depth of the network. As a result, the model is better able to capture fine-grained features necessary for distinguishing between frequent and rare labels. Furthermore, the depth of ResNet50 enables it to build highly abstract and hierarchical representations, making it well-suited for complex multi-label datasets where different labels may exhibit varying levels of visual similarity. Another important factor is ResNet50's strong generalization capability. In the context of a long-tailed distribution, models tend to overfit on dominant classes while neglecting rare ones. ResNet50, owing to its design, offers better regularization and smoother convergence during training, helping the model maintain balanced performance across classes. For these reasons, ResNet50 stands out as a strong candidate for handling the long-tailed, multi-label classification challenges presented by our dataset.

4.2.2 Architectural Overview of ResNet50

ResNet50 is a deep convolutional neural network composed of 50 layers, structured around the concept of residual blocks. Each residual block includes a set of convolutional layers with shortcut connections that add the input directly to the output of a few stacked layers. This structure allows the network to maintain the flow of gradients during backpropagation, which is essential for training deeper architectures without suffering from degradation problems.

In the context of our long-tailed, multi-label dataset, the architecture of ResNet50 offers several advantages. The initial layers of ResNet50 focus on extracting low-level features such as edges and textures, which are crucial for all classes, including minority classes. As the network deepens, it captures higher-level, more abstract features that are vital for distinguishing complex patterns associated with specific labels. The bottleneck design — where three convolutional layers (1x1, 3x3, 1x1) form a single residual block — ensures computational efficiency while maintaining high representational power.

Moreover, ResNet50's global average pooling layer before the final fully connected layer prevents the model from heavily relying on spatial information, making it more adaptable to intra-class variations. This is particularly beneficial for rare classes, where examples may be visually diverse yet semantically similar. In our experiments, we observed that ResNet50's architectural features significantly improved the model's ability to generalize across both head and tail classes, resulting in better balanced performance across the label spectrum.




Fig 4.1 Architecture of ResNet50

The architecture of ResNet50 featuring deep residual learning through skip connections. The model comprises convolutional layers organized into bottleneck residual blocks, enabling efficient training of very deep networks by mitigating vanishing gradient issues.

4.2.3 InceptionV3

InceptionV3 is a convolutional neural network architecture renowned for its efficiency and accuracy in image classification tasks. Its design incorporates factorized convolutions, auxiliary classifiers, and batch normalization, which collectively enhance its performance and reduce computational complexity.

In the context of long-tailed multi-label classification, InceptionV3's ability to capture multi-scale features is particularly beneficial. The architecture's inception modules process input data through multiple convolutional filters of varying sizes simultaneously, enabling the model to learn both fine-grained and coarse features. This multi-scale feature extraction is crucial for recognizing classes with limited samples, as it allows the model to generalize better from sparse data.

Moreover, the inclusion of auxiliary classifiers during training provides additional gradient signals, which help in mitigating the vanishing gradient problem in deep networks. This feature ensures that even the earlier layers in the network learn meaningful representations, which is essential when dealing with underrepresented classes in a long-tailed distribution.

4.2.4 Architectural Overview of InceptionV3

InceptionV3's architecture is composed of multiple inception modules, each designed to perform convolutions with different filter sizes (1x1, 3x3, 5x5) and pooling operations in parallel. This design allows the network to capture diverse spatial features at each level. The use of 1x1 convolutions before the larger convolutions serves as a dimensionality reduction step, decreasing computational cost while preserving essential information.

The network begins with a series of convolutional and pooling layers that reduce the spatial dimensions of the input image. This is followed by a sequence of inception modules that extract complex features. The final layers consist of global average pooling, dropout for regularization, and a fully connected layer that outputs the class probabilities.

For your dataset, which likely contains images with varying object sizes and complex backgrounds, InceptionV3's architecture is well-suited. Its ability to process information at multiple scales ensures that both prominent and subtle features are learned effectively, aiding in the accurate classification of both frequent and rare labels.




Fig 4.2 : Architecture of InceptionV3

The architecture of InceptionV3 built around inception modules that perform parallel convolutions with different kernel sizes. This multi-branch structure enables the model to effectively capture features at multiple scales, improving classification performance across diverse and complex datasets.

4.2.5 VGG16

VGG16 is a deep convolutional neural network known for its simplicity and uniform architecture. It consists of 13 convolutional layers followed by 3 fully connected layers, using small 3x3 filters throughout the network. This consistent structure allows for straightforward implementation and has been effective in various image classification tasks.

In the realm of long-tailed multi-label classification, VGG16's depth enables it to learn hierarchical feature representations, which can be advantageous for distinguishing between classes. However, its lack of specialized mechanisms to handle class imbalance means that, without additional techniques such as data augmentation or loss function adjustments, it may struggle to accurately classify underrepresented classes.

4.2.6 Architectural Overview of VGG16

The VGG16 architecture follows a sequential pattern, stacking multiple convolutional layers with ReLU activation functions, interspersed with max-pooling layers to reduce spatial dimensions. After the convolutional blocks, the network includes three fully connected layers, culminating in a softmax output layer.

For your dataset, VGG16 can serve as a strong baseline model. Its straightforward design allows for easy modifications, such as incorporating class weighting or focal loss to address class imbalance. Additionally, its pre-trained weights on large datasets like ImageNet can be fine-tuned to adapt to the specific characteristics of your data, potentially proving performance on minority classes.




Fig 4.3: Architecture of VGG16

The architecture of VGG16 characterized by a straightforward, uniform design of stacked 3x3 convolutional layers and max pooling layers. The network transitions into fully connected layers for classification, facilitating effective hierarchical feature extraction.

4.2.7 ResNet101

ResNet101 extends the principles of the ResNet architecture by increasing the depth to 101 layers, incorporating residual connections that facilitate the training of very deep networks. These connections allow gradients to flow more easily during backpropagation, mitigating issues like vanishing gradients.

In the context of long-tailed multi-label classification, ResNet101's depth enables it to learn complex and abstract features, which can be particularly beneficial for distinguishing between classes with subtle differences. The residual connections ensure that learning is efficient and that the network can generalize well, even when some classes have limited training samples.

4.2.8 Architectural Overview of ResNet101

ResNet101 is composed of an initial convolutional layer followed by multiple residual blocks, each containing convolutional layers and identity shortcuts. These blocks allow the network to learn residual functions, effectively enabling the training of deeper networks without degradation in performance.

For your dataset, ResNet101's architecture is advantageous due to its capacity to model intricate

patterns and relationships within the data. Its depth allows for the extraction of high-level features necessary for accurate multi-label classification, while the residual connections ensure that learning remains efficient and effective across all classes, including those that are underrepresented.




Fig 4.4: Architecture of ResNet101

The architecture of ResNet101 extending the residual learning framework with a greater number of layers. The deep network structure enables the extraction of highly abstract features, supporting robust learning even in the presence of class imbalance.

4.3 Loss Functions Experimented

One of the core challenges in multi-label classification, especially in the medical imaging domain, is handling **class imbalance** and **co-occurrence of labels**. In the ODIR-5K dataset, some ocular disease categories were significantly under-represented compared to others, leading to biased training if a standard loss function is used. To address this, a variety of **loss functions** were explored, each tailored to the specific nuances of the dataset and the multi-label nature of the problem.

- **Binary Cross Entropy (BCE):**

This is the default loss function for multi-label classification problems. It treats each class as a separate binary classification task, making it well-suited for situations where multiple classes can be present simultaneously. BCE calculates the average of the binary loss across all labels, which makes it computationally efficient and easy to implement. However, it does not inherently account for class imbalance, which can hinder performance for rare disease categories.

- **Weighted Binary Cross Entropy (Weighted BCE):**

To mitigate the limitations of standard BCE, class weighting was introduced. This method assigns higher weights to minority classes during training, forcing the model to focus more on under-represented labels. Weighted BCE improved the overall recall and precision for rare diseases without significantly affecting performance on the majority

classes. The class weights were computed inversely proportional to label frequency, a common practice in imbalanced learning.

- **Hill Loss:**

Hill Loss is particularly designed for long-tailed and imbalanced datasets, where few-shot classes coexist with many-shot ones. It dynamically adjusts the penalty in the loss function by modifying the shape of the loss curve, leading to greater gradient feedback for rare classes. This allows the model to learn better representations for difficult samples and minority labels. Although it did not outperform other advanced techniques in isolation, Hill Loss contributed positively when used in hybrid combinations.

- **Dice Loss:**

Initially developed for semantic segmentation tasks, Dice Loss measures the overlap between predicted labels and ground truth labels. In multi-label classification, it helps to directly optimize for F1-score-like behaviour, which is critical when dealing with imbalanced datasets. Dice Loss gave a noticeable improvement in the macro-averaged performance, especially where standard BCE struggled.

- **Hybrid Losses (BCE + Hill, Dice + Hill):**

Recognizing that no single loss function was optimal across all label categories, hybrid losses were formulated to combine their strengths. BCE + Hill provided a balance between the stability of BCE and the class sensitivity of Hill Loss, while Dice + Hill further improved recall and F1 scores, especially for rare conditions. These combinations yielded the most consistent results, with Dice + Hill emerging as the best-performing configuration during validation

Loss Function	Validation Accuracy (%)	Observations
BCE	58.4	Baseline performance; lacks imbalance handling
Weighted BCE	61.2	Better performance on under-represented classes
Hill Loss	60.5	Stronger performance on rare classes
Dice Loss	62.3	Improved overlap and F1-score
BCE + Hill	64.1	Balanced accuracy and stability
Dice + Hill	65.7	Best overall; high recall and robustness

These findings underscore the importance of loss function engineering in medical imaging tasks. Selecting the right objective not only accelerates convergence but significantly improves the quality of predictions, especially in datasets with skewed label distributions.

4.4 Model Training Strategy

The training methodology was meticulously designed to achieve optimal model convergence, minimize overfitting, and enhance generalization to unseen data. Each step of the training pipeline was carefully considered to ensure both robustness and reproducibility. The following strategies were employed:

- **Data Splitting:**
A **Stratified 5-Fold Cross-Validation** approach was used to address the dataset's class imbalance. This ensured that each fold maintained the original label distribution, providing a more reliable and unbiased estimation of model performance across multiple trials.
- **Data Augmentation:**
Basic augmentation techniques were applied, including random horizontal and vertical flips, slight image rotations, and translations. While more advanced methods like brightness/contrast adjustments, cutout, and random erasing were not incorporated at this stage, the existing augmentations significantly helped in reducing the risk of overfitting by increasing dataset diversity.
- **Fine-Tuning and Trainable Layers:**
All models were initialized with **completely unfrozen layers** from the beginning, enabling full fine-tuning of pretrained weights. In certain trials, **layer-wise unfreezing** was strategically employed — where only deeper layers were unfrozen after five epochs — to allow better low-level feature retention while adapting higher-level features to the specific task.
- **Regularization Techniques:**
Dropout layers with rates varying between **0.25 and 0.4** were incorporated depending on the model's depth and complexity. This regularization helped prevent overfitting by randomly deactivating neurons during training, thereby promoting better generalization.
- **Optimization Strategy:**
The **Adam optimizer** was chosen for its adaptive learning rate capability, initialized at **1e-4**. Furthermore, **learning rate decay** and Adam Optimizer was employed, dynamically lowering the learning rate when validation metrics plateaued, allowing finer convergence without abrupt updates.
- **Training Environment and Reproducibility:**
All training experiments were conducted on **Google Colab** with **GPU acceleration** (typically Tesla T4 or P100 GPUs). A clean, modular training pipeline was implemented, ensuring that experiments could be easily reproduced, parameters systematically varied, and results accurately recorded for comparative analysis.

4.5 Evaluation Metrics

To comprehensively evaluate the performance of the multi-label classification models, several advanced metrics were adopted that go beyond traditional top-1 accuracy. Given the complexity of the ODIR-5K dataset—where each image may belong to multiple ocular disease categories—it was imperative to use metrics that reflect both individual label prediction and the overall predictive balance of the model.

- **Threshold Multi-Label Accuracy:**

Accuracy was initially measured by applying a threshold of 0.5 on the sigmoid outputs. This converts each predicted probability into a binary decision per class. However, this default threshold might not yield the best performance across all classes. Some labels may require a more relaxed or stringent threshold depending on their prevalence. To address this, further optimization using ROC-AUC-based threshold tuning is planned as future work. This will allow setting class-specific thresholds that better separate positive from negative predictions.

- **Precision, Recall, and F1-score:**

These three metrics were computed per class and then averaged using macro (unweighted average) and micro (global) methods. This approach ensures the model's performance is evaluated both from the perspective of individual diseases (important for rare conditions like Hypertension or Myopia) and overall classification consistency. Precision indicates the model's ability to avoid false positives, recall reflects its capability to detect true positives, and F1-score provides a harmonic mean of the two. These metrics are particularly critical in medical diagnostics where both overdiagnosis and underdiagnosis can have serious implications.

- **Confusion Matrix (Adapted):**

Although multi-label problems do not traditionally use confusion matrices, we adapted this visualization by flattening predictions into binary vectors. This allowed us to observe which classes frequently co-occurred or were misclassified together. For example, Diabetic Retinopathy and Cataract occasionally overlapped in prediction errors due to similar visual artifacts in fundus images.

- **ROC-AUC (To Be Integrated):**

The Receiver Operating Characteristic - Area Under Curve is a standard evaluation metric in medical machine learning due to its effectiveness in reflecting model robustness across all classification thresholds. Although not incorporated into the current version of the model, we aim to utilize per-class ROC-AUC scores in the next iteration. These scores will help identify how well each class is distinguishable from others and will inform the threshold tuning process.

The figure below provides a comparative visualization of **Precision, Recall, and F1-score** for the four main architectures evaluated in this work. It is evident that ResNet101 provides a

balanced improvement across all three metrics, showcasing its superior capability to generalize across both common and rare ocular diseases.




Fig 4.5 Comparison Graph of different Models

After extensive experimentation, ResNet101 emerged as the best-performing model with a test accuracy of 66%, leveraging deeper feature extraction capabilities. ResNet50 followed closely with 63.1% accuracy, showing stable and fast convergence when trained with the Dice+Hill hybrid loss. InceptionV3 achieved a balanced performance with 57.8% accuracy using Dice Loss, while VGG16, being a simpler and quicker-to-train model, reached 55.2% accuracy with Weighted Binary Cross-Entropy. Overall, the combination of advanced architectures and the use of hybrid loss functions like Dice+Hill significantly improved model performance.

CHAPTER - 5

ADVANCED TECHNIQUES FOR ROBUSTNESS AND ACCURACY ENHANCEMENT

5.1 GENERAL

In previous experiments, we observed that using a single deep learning model—such as ResNet50, VGG16, InceptionV3, or ResNet101—in combination with a single loss function, such as Binary Cross-Entropy (BCE) loss or Weighted Binary Cross-Entropy loss, did not lead to a significant improvement in the accuracy of the parameter under investigation. The marginal gains suggested that relying on a single loss function might not be sufficient to fully capture the complexities of the learning task.

Consequently, we shifted our approach toward exploring hybrid loss functions, which aim to leverage the complementary strengths of multiple loss terms. For instance, a hybrid loss combining Binary Cross-Entropy (BCE) loss and Hinge loss offers several advantages. BCE loss is effective for probabilistic binary classification, focusing on minimizing the difference between predicted probabilities and true labels. However, it does not explicitly enforce a strong margin between classes. In contrast, Hinge loss, traditionally used in support vector machines, encourages a larger margin between classes, thereby promoting better class separation and generalization.

By combining BCE loss and Hinge loss into a hybrid loss function, the model benefits from both probabilistic calibration (through BCE) and robust class separation (through Hinge loss). This dual optimization strategy allows the model not only to predict the correct class with high probability but also to maintain a safe margin between decision boundaries, thereby improving both accuracy and robustness, particularly in imbalanced or noisy datasets.

Thus, the use of hybrid loss functions presents a more nuanced optimization landscape, leading to superior performance compared to using any individual loss function alone.

5.2 LITERATURE REVIEW

Loss functions are at the heart of training deep learning models as they provide a measure of how well a model's predictions match the true labels. These functions guide the optimization process, driving model learning. In early stages of experimentation, standard loss functions such as Binary Cross-Entropy (BCE) loss and Weighted Binary Cross-Entropy (Weighted BCE) loss were applied to popular architectures, including VGG16, InceptionV3, ResNet50, and ResNet101. However,

these models exhibited only marginal improvements in performance, especially when dealing with complex and imbalanced datasets, which are common in practical applications.

This observation led to a more thorough investigation into alternative and hybrid loss formulations. The goal was to combine the strengths of multiple loss functions to better address the challenges posed by datasets with a long-tailed distribution, particularly in multi-label classification tasks. Long-tailed distributions refer to datasets where a small number of classes are heavily represented, while the majority of classes have very few instances. In such cases, the model may become biased towards the majority class and underperform on the minority classes.

In this section, we explore key loss functions, their mathematical foundations, and how they are relevant in the context of long-tailed distributions in multi-label datasets. We also present findings from relevant literature that inspired their use.

5.2.1 Binary Cross-Entropy (BCE) Loss

Binary Cross-Entropy (BCE) loss, also known as log loss, is one of the most widely used loss functions in binary classification tasks. It measures the dissimilarity between the predicted probabilities of the model and the actual binary labels. BCE loss is particularly well-suited for tasks where the model's output is a probability between 0 and 1.

The formula for BCE loss is:

$$BCE\ Loss = -\frac{1}{N} \sum_{i=1}^N [\alpha y_i \log(p_i) + (1 - y_i) \log(1 - p_i)] \quad (1)$$

Where α is the weight for the positive class, typically greater than 0.5 when the positive class is rare.

By incorporating a weight factor α , this loss function ensures that the minority class is given more attention during training, helping to mitigate the effects of class imbalance. This is particularly useful in multi-label tasks with long-tailed distributions, where the model might otherwise learn to predict the majority class at the expense of the minority classes.

5.2.2 Weighted Binary Cross-Entropy (Weighted BCE) Loss

To address the issue of class imbalance, Weighted BCE loss introduces a weighting factor, α , which assigns greater importance to the minority class. This ensures that the model does not ignore the underrepresented classes in long-tailed datasets. The formula for Weighted BCE loss is:

$$Weighted\ BCE\ Loss = -\frac{1}{N} \sum_{i=1}^N [\alpha y_i \log(p_i) + (1 - \alpha)(1 - y_i) \log(1 - p_i)] \quad (2)$$

Where α is the weight for the positive class, typically greater than 0.5 when the positive class is rare.

By incorporating a weight factor α , this loss function ensures that the minority class is given more attention during training, helping to mitigate the effects of class imbalance. This is particularly useful in multi-label tasks with long-tailed distributions, where the model might otherwise learn to predict the majority class at the expense of the minority classes.

5.2.3 Hill Loss

Hill loss is designed to combine the benefits of smooth probabilistic predictions with margin enforcement, similar to Hinge loss. The key idea behind Hill loss is that it encourages not only correct predictions but also confident predictions by enforcing a margin between the predicted probabilities and the true labels. This helps the model make clearer, more confident decisions.

The mathematical formulation for Hill loss is:

$$\text{Hill Loss}(p, y) = \max(0, m - (2y - 1)p)^2 \quad (3)$$

Where:

- m is the margin hyperparameter (typically $m=1$),
- p is the predicted probability,
- y is the true label.

Hill loss helps in scenarios where the model faces noisy labels or ambiguous samples. In the case of long-tailed distributions, Hill loss can be beneficial because it encourages the model to become more confident in its predictions, reducing the tendency to give uncertain predictions, especially for minority classes. The margin enforcement aspect also ensures that the model doesn't overly favor the majority class, thus improving performance on underrepresented classes.

5.2.4 Dice Loss

Dice loss is often used in segmentation tasks but can be adapted for binary classification to handle severe class imbalance by focusing on the overlap between predicted and ground truth areas. It measures the similarity between two sets and is defined by the Dice coefficient, which is particularly effective in imbalanced datasets.

The Dice coefficient is given by:

$$\text{Dice Coefficient} = \frac{2 \sum_{i=1}^N p_i y_i}{\sum_{i=1}^N p_i + \sum_{i=1}^N y_i} \quad (4)$$

Thus, the Dice loss can be expressed as:

$$\text{Dice Loss} = 1 - \text{Dice Coefficient} \quad (5)$$

Dice loss is highly relevant for multi-label classification tasks with long-tailed distributions, as it prioritizes the overlap between true positives and predicted positives. In such datasets, where the majority class often dominates, Dice loss ensures that the model does not overlook smaller, underrepresented classes by focusing on the quality of overlap. This makes it a powerful tool for handling class imbalance in both segmentation and classification tasks.

5.2.5 Hybrid BCE + Hill Loss

The Hybrid BCE + Hill Loss combines the probabilistic prediction capabilities of BCE loss with the margin enforcement properties of Hill loss. This hybrid approach helps to address both the class imbalance and the need for confident predictions in challenging scenarios like long-tailed distributions.

The hybrid loss function is:

$$\text{Hybrid BCE} + \text{Hill Loss} = \lambda_1 \times \text{BCE Loss} + \lambda_2 \times \text{Hill Loss} \quad (6)$$

For long-tailed datasets, this hybrid loss is highly beneficial as it combines the strengths of both Dice loss (which handles imbalances) and Hill loss (which enforces confidence). The model is thus encouraged to not only focus on class overlap but also on learning more confident and robust predictions, especially for minority classes that are underrepresented in the data.

5.3 CONCLUSIONS

In the context of multi-label classification tasks with long-tailed distributions, each loss function evaluated offers distinct advantages, but also limitations depending on the dataset and model architecture. Here's a summary of how each loss function fares:

- Binary Cross-Entropy (BCE) Loss is the simplest and most widely used, but struggles in imbalanced datasets. Its tendency to favor the majority class often results in poor

performance on minority classes, making it less effective in handling long-tailed distributions.

- Weighted Binary Cross-Entropy (Weighted BCE) Loss addresses this imbalance by introducing a weighting factor, α (alpha), to give more importance to the minority class. This is a significant improvement over standard BCE for long-tailed datasets, but it still relies on the assumption that the class imbalance can be captured solely by the weights, which might not be sufficient in more complex distributions.
- Hill Loss introduces the concept of margin-based confidence enforcement, which helps the model make more confident decisions and avoid ambiguity. This is particularly beneficial in scenarios where the dataset is noisy or uncertain, but its effect on long-tailed distributions is less direct than other methods.
- Dice Loss shines in highly imbalanced settings, particularly when there is a need to focus on the overlap between predicted and true labels, which makes it very relevant for long-tailed distributions. It provides a better measure of performance when the model needs to ensure it doesn't ignore minority classes, though it may still be sensitive to certain biases in the data.
- Hybrid BCE + Hill Loss combines the strengths of BCE loss (probabilistic predictions) with Hill loss (confidence enforcement), offering a more robust approach to handling long-tailed distributions. It works well for multi-label tasks where both class imbalance and uncertain predictions are prevalent. This hybrid approach tends to outperform simpler loss functions due to its ability to manage both imbalance and prediction quality.
- Hybrid Dice + Hill Loss takes the strengths of Dice loss (overlap focus) and Hill loss (confidence enforcement) to a new level. This hybrid loss is highly beneficial for long-tailed datasets, particularly in scenarios where both precise class overlap and model confidence are critical.




Fig 5.1: Performance comparison of different neural network architectures using multiple loss functions on the ODIR-5K dataset. The figure illustrates the variation in accuracy

achieved across BCE, Weighted BCE, Hill Loss, Dice Loss, and hybrid loss formulations, highlighting the superior performance of ResNet-based models under long-tailed multi-label conditions.

In conclusion, the Hybrid Dice + Hill Loss emerges as the most effective choice for handling long-tailed distributions in multi-label classification tasks, as it combines both the class overlap focus of Dice loss and the confidence enforcement of Hill loss. Weighted BCE Loss and Hybrid BCE + Hill Loss are also strong contenders, especially when dealing with imbalance but without compromising the model's confidence in its predictions. While BCE Loss and Hill Loss offer certain benefits, they are generally less effective on their own when handling long-tailed data. Dice Loss provides a focused solution but may require further refinement in more complex imbalances.

Architecture	BCE Acc	Weighted BCE Acc	Hill Loss Acc	Dice Loss Acc	Hybrid BCE+Hill Acc	Hybrid Dice+Hill Acc	Trainable Params	Dropout Rate
VGG16	55%	61%	52%	54%	58%	56%	~15 Million	0.3
InceptionV3	54%	56%	52%	53%	57%	58%	~23 Million	0.4
ResNet50	59%	61%	57%	58%	62%	63%	~25 Million	0.3
ResNet101	66%	67%	64%	64%	62%	63%	~44 Million	0.25

Ultimately, the choice of loss function should be guided by the nature of the dataset and the model architecture. For datasets with severe class imbalance, hybrid approaches like Hybrid Dice + Hill Loss should be considered as they offer a comprehensive solution to both class imbalance and prediction confidence.

CHAPTER 6

RESULTS AND CONCLUSION

6.1 RESULT

6.1.1 Overall Performance

The four deep convolutional neural network backbones VGG16, InceptionV3, ResNet50, and ResNet101 were evaluated on the held out test set of the ODIR-5K dataset using a stratified five-fold cross-validation framework. This strategy ensured that the inherent long-tail distribution of classes was preserved across training and validation splits, providing a robust measure of generalization.

The final performance results are summarized in Table 6.1. Each model was trained with its optimal loss configuration. Among the models tested, ResNet101 trained with the weighted Binary Cross-Entropy (BCE) achieved the highest test accuracy of **66.0%** and a macro-averaged F1-score of **0.62**. ResNet50, also using the Dice + Hill hybrid loss, attained an accuracy of **63.1%** and a macro F1-score of **0.59**. InceptionV3, optimized with the Dice loss, reached a test accuracy of **58.0%**, whereas VGG16 trained with Weighted BCE achieved **55.2%** accuracy.

These results emphasize that deeper architectures and hybrid loss functions substantially contribute to the improved detection of multiple ocular disease labels in fundus images.

Model	Loss	Test Accuracy	Macro F1
VGG16	Weighted BCE	55.2%	0.48
InceptionV3	Dice	58.0%	0.51
ResNet50	Dice + Hill	63.1%	0.59
ResNet101	Weighted BCE	66.0%	0.62

Table 6.1: Test Performance by Model and Loss

6.1.2 Loss Function Comparison

A detailed analysis of the impact of different loss functions revealed that hybrid loss combinations consistently outperformed single losses across all backbones. When trained with unweighted Binary Cross-Entropy (BCE), the models achieved a baseline test accuracy of approximately **58%**, but performance on minority labels was notably poor, reflecting the severe class imbalance.

Applying Dice + Hill increased the overall accuracy to around **61%**, demonstrating that rebalancing the class contributions during training improves minority class sensitivity. The use of Dice loss further elevated the performance, achieving roughly **62%** accuracy by explicitly maximizing overlap between predicted and ground-truth labels. Training with Hill loss alone led to approximately **60.5%** accuracy, emphasizing better correction of false negatives but requiring careful tuning of margin parameters.

Among all tested configurations, the hybrid combinations of losses yielded the highest improvements. BCE + Hill loss achieved **64.1%** accuracy, balancing the model's stability and its ability to detect rare conditions. However, the best-performing loss was Weighted BCE, which achieved **66.0%** accuracy, effectively combining the strengths of overlap-based optimization and margin-based robustness to class imbalance.




Fig 6.1 Comparison of Architecture using Weighted BCE Loss

The superior performance of hybrid losses validates the hypothesis that addressing both class overlap and positive instance suppression simultaneously leads to better multi-label prediction performance in long-tail distributions.

6.1.3 Cross-Validation Stability

The stability of model performance across folds was a critical consideration in evaluating the robustness of the training protocol. Using stratified five-fold cross-validation, each model was retrained on four folds and tested on the remaining fold, repeating the process for all five splits.

For the ResNet101 model trained with the Dice + Hill loss, the mean accuracy across the five folds was consistently around **66.0%**, with a low standard deviation of $\pm 2.3\%$. This small variance across folds highlights the model's generalization capacity, suggesting that the training procedure and data augmentation strategies were successful in mitigating overfitting, even under extreme class imbalance.

Such stable cross-validation performance is crucial for multi-label medical imaging applications, where generalization to unseen samples is essential for clinical reliability.

6.1.4 Summary of Key Results

Aspect	Key Result
Effect of Hybrid Loss	Dice + Hill hybrid loss improved accuracy from 58.4 % (BCE baseline) to 66.0 %—an absolute gain of 7.6 %.
Cross-Validation Stability	Stratified 5-fold CV yielded mean accuracy = 66.0 %, with a low standard deviation of ± 2.3 %.
ROC-AUC	Average per-class ROC-AUC ≈ 0.84 , indicating strong discriminative power across all eight labels.
Data Augmentation Impact	Spatial augmentations (rotations, flips, zoom) increased recall on minority classes by ≈ 5 %.
Model Calibration	Sigmoid outputs showed slight over-confidence on majority classes; temperature scaling recommended.
Backbone Depth Advantage	ResNet101 (101 layers) outperformed ResNet50 (+2.9 %), InceptionV3 (+8.0 %), and VGG16 (+10.8 %).

6.2 Conclusion

6.2.1 Key Findings

The results presented in this work demonstrate that a carefully designed deep learning pipeline can significantly improve multi-label classification performance on long-tail medical datasets. Specifically, the depth of the ResNet101 backbone proved critical for capturing the complex and subtle features associated with multiple ocular pathologies, outperforming shallower models by a notable margin.

Moreover, hybrid loss functions, particularly the Dice + Hill combination, consistently delivered superior results compared to standard single-term losses. This confirms that simultaneously optimizing for overlap and positive label preservation addresses both the negative dominance problem and the extreme class imbalance inherent in multi-label medical data.

Finally, the robust performance observed across stratified folds affirms that the model's generalization capacity is strong, an essential prerequisite for clinical applicability.

6.2.2 Implications

This study establishes a validated and reproducible pipeline for multi-label ocular disease classification using deep learning. It offers a practical framework that integrates advanced backbone architectures, hybrid loss functions, and stratified validation to address the specific challenges posed by long-tail distributions. The methods and findings here can serve as a strong baseline for future work in medical image classification tasks where multiple pathologies coexist.

6.2.3 Limitations

Despite promising results, several limitations remain. The ODIR-5K dataset is geographically restricted, potentially limiting the generalizability of the models to populations with different ocular disease prevalence and imaging characteristics. Additionally, the computational demands of training deep residual networks with hybrid losses are non-trivial, posing challenges for real-time deployment or deployment on edge devices without further optimization.

6.2.4 Future Work

Future research directions include the implementation of per-class threshold tuning based on ROC-AUC curves to maximize clinical utility, especially for rare disease classes. Integrating few-shot and zero-shot learning methods could enable the model to generalize to unseen classes without requiring extensive retraining. Finally, model compression techniques, including pruning and quantization, will be investigated to facilitate efficient deployment in resource-constrained environments.

References

1. Kingma, D. P., & Ba, J. (2015). *Adam: A Method for Stochastic Optimization*. arXiv:1412.6980. <https://arxiv.org/abs/1412.6980>
2. Bertinetto, L., Henriques, J. F., Valmadre, J., Torr, P. H. S., & Vedaldi, A. (2020). *Making Better Mistakes: Leveraging Class Hierarchies with Deep Networks*. arXiv:2006.07839. <https://arxiv.org/abs/2006.07839>
3. Milletari, F., Navab, N., & Ahmadi, S. A. (2016). *V-Net: Fully Convolutional Neural Networks for Volumetric Medical Image Segmentation*. Proceedings of the Fourth International Conference on 3D Vision (3DV), 565–571. <https://doi.org/10.1109/3DV.2016.79>
4. Wu, T., Tang, Y., Wang, Y., Hu, R., & Lin, J. (2021). *Robust Asymmetric Loss for Multi-Label Long-Tailed Learning*. Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), 10964–10973. <https://doi.org/10.1109/ICCV48922.2021.01079>
5. Chen, S., Liu, Y., Zha, X., & Yin, B. (2020). *Multi-Label Classification of Fundus Images With Graph Convolutional Network and Self-Supervised Learning*. Frontiers in Cell and Developmental Biology, 8, 595831. <https://doi.org/10.3389/fcell.2020.595831>
6. He, K., Zhang, X., Ren, S., & Sun, J. (2016). *Deep Residual Learning for Image Recognition*. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 770–778. <https://doi.org/10.1109/CVPR.2016.90>
7. Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., & Wojna, Z. (2016). *Rethinking the Inception Architecture for Computer Vision*. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2818–2826. <https://doi.org/10.1109/CVPR.2016.308>
8. Deep Learning-Based Pipeline for the Detection of Multiple Ocular Diseases, Machine Learning and Computational Intelligence Techniques for Data Engineering (MISP 2022), May 16, 2023, DOI: 10.1007/978-981-99-0047-3_69.
9. Pre-Trained Deep Learning-Based Approaches for Eye Disease Detection, Pankaj Kumar, Sudhir Bhandari, Vishal Dutt, 2023 International Conference on Circuit Power and Computing Technologies (ICCPCT), August 2023, DOI: 10.1109/ICCPCT58313.2023.10245175.
10. An Efficient Deep Learning Model for Eye Disease Classification, Archana Saini, Kalpna Guleria, Shagun Sharma, 2023 International Research Conference on Smart Computing and Systems Engineering (SCSE), June 29, 2023, DOI: 10.1109/SCSE59836.2023.10215000.
11. Explainable Classification of Macular Degeneration Using Deep Learning, Sahil Khose, Ankita Ghosh, Yogish S. Kamath, Neetha I.R. Kuzhuppilly, J. R. Harish Kumar, Manipal Academy of Higher Education, published online at researcher.manipal.edu.

BIO DATA

	Name: Prashant Mishra Date of Birth: 05/12/2003 Education: Undergrad in Electrical Engineering Email: prashant66m@gmail.com Phone: +91 9140637106
	Name: K Abhishek Raj Date of Birth: 28/07/2003 Education: Undergrad in Electrical Engineering Email: abhishek28k7@gmail.com Phone: +91 7893051444
	Name: Rakesh Yamjala Date of Birth: 16/04/2004 Education: Undergrad in Electrical Engineering Email: rakeshyamjala04@gmail.com Phone: +91 6281931668
	Name: Sandeep K Aazad Date of Birth: 01/03/2003 Education: Undergrad in Electrical Engineering Email: sandeepkumaraazad622@gmail.com Phone: +91 9140536637